

A SURVEY OF INTERNET WORM DETECTION AND CONTAINMENT

PELE LI, MEHDI SALOUR, AND XIAO SU, SAN JOSE STATE UNIVERSITY

ABSTRACT

Self-duplicating, self-propagating malicious codes known as computer worms spread themselves without any human interaction and launch the most destructive attacks against computer networks. At the same time, being fully automated makes their behavior repetitious and predictable. This article presents a survey and comparison of Internet worm detection and containment schemes. We first identify worm characteristics through their behavior, and then classify worm detection algorithms based on the parameters used in the algorithms. Furthermore, we analyze and compare different detection algorithms with reference to the worm characteristics by identifying the type of worms that can and cannot be detected by these schemes. After detecting the existence of worms, the next step is to contain them. This article explores the current methods used to slow down or stop the spread of worms. The locations to implement detection and containment, as well as the scope of each of these systems/methods, are also explored in depth. Finally, this article points out the remaining challenges of worm detection and future research directions.

Self-propagating malicious codes known as computer worms spread themselves without any human interaction and launch the most destructive attacks against computer networks. Being fully automated, a worm's behavior is usually repetitious and predictable, making it possible to be detected.

A worm's life consists of the following phases: target finding, transferring, activation, and infection. Since worms involve network activities in the first two phases, their behaviors in these two phases are critical for developing detection algorithms. Therefore, this article first focuses on worm characteristics that facilitate their detection.

Many algorithms have been proposed in the past to try to catch and stop the spread of Internet worms. Most research papers discuss efforts that are related to their proposed work, but none of these papers gives a comprehensive classification of the existing detection and containment systems. This article contains a survey and analysis of Internet worm detection and containment systems. Our research categorizes these systems based on the parameters used in each scheme. These categories are compared against worm characteristics, and the insufficiency of current systems is pointed out.

After detecting the existence of worms, the next step is to contain them. This article explores the current methods used to slow down or stop the spread of worms. The locations to

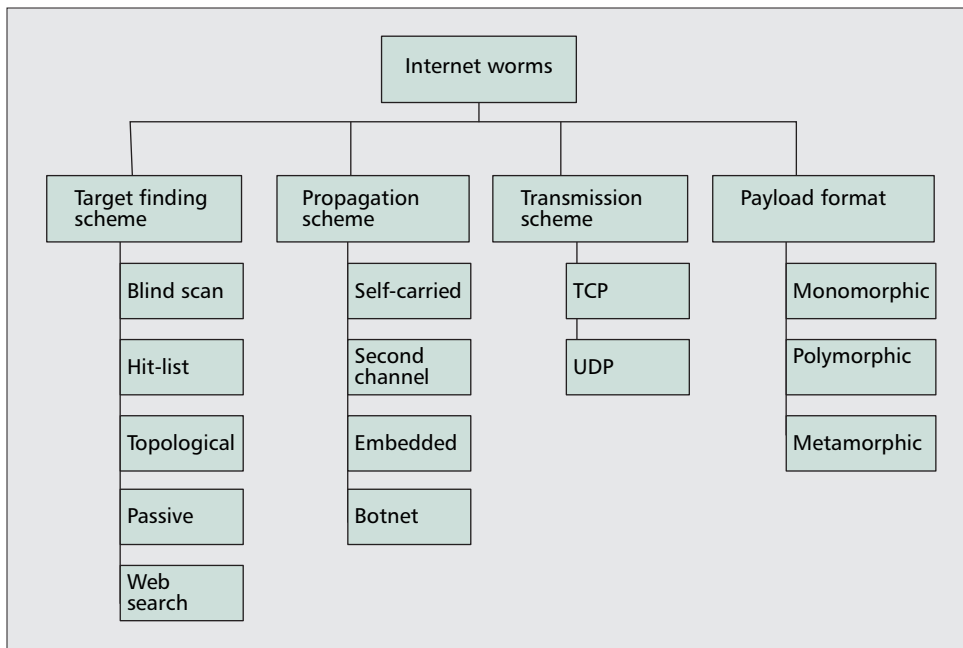
implement detection and containment, as well as each of these system scopes, are also explored in depth at each level.

OVERVIEW

After an introductory terminology section is presented, worm characteristics during target finding and worm transferring phases are identified. This is followed by an overview of worm defense mechanisms: detection and containment. The classification of detection algorithms and containment systems is presented next. Depending on where the detection and containment systems are implemented, they may construct different views of worm propagation behaviors, so there may be differences in the scope of their defenses. We conclude the survey by identifying future research challenges.

TERMINOLOGY

- *Activation*: Activation is when a worm starts performing its malicious activities. Activation might be triggered on a specific date or under certain conditions.
- *False alarm*: A false alarm is an incorrect alert generated by a worm detection system.
- *False positive*: A false positive is a false alarm where an alert is generated when there is no actual attack or



■ **Figure 1.** Categorization of worm characteristics.

threat.

- **False negative:** False negative means the detection system missed an attack. It is a false negative if no alert is generated while the system is under an attack.
- **Infection:** Infection is the result of the worm performing its malicious activities on the host.
- **Target finding:** Target finding is the first step in a worm's life to discover victims (vulnerable hosts).
- **Threshold:** Threshold is a predefined condition that, if met, indicates the existence of specious traffic or a worm attack.
- **Transfer:** Transfer refers to sending a copy of the worm to the target after the victim (target) is discovered.
- **Virus:** A virus is a malicious piece of code that attaches to other programs to propagate. It cannot propagate by itself, and normally depends on a certain user intervention, such as opening up an email attachment or running an executable file, to be activated [1].
- **Worm:** A worm is a malicious piece of code that self-propagates, often via network connections, exploiting security flaws in computers on the network. In general, worms do not need any human intervention to propagate; however, a category of worms called *passive* worms require certain host behavior or human intervention to propagate. For example, a passive worm only propagates itself until it is contacted by another host.

INTERNET WORMS

Since the Morris worm in 1988, Internet worms have caused the most extensive and widespread damage of all kinds of computer attacks. An Internet worm is defined as a piece of malicious code that duplicates and propagates by itself. Usually, it does not require any human interaction and spreads via network connections.

The life of a worm, after it is released, typically includes the following phases: target finding, worm transferring, worm activation, and infection. During the phase of target finding and worm transferring, the worm is active over the Internet, making it possible for network-based intrusion detection systems (NIDSs) to catch the worm. The activities in the two latter

phases are limited to local machines and are harder to detect by NIDSs. This article categorizes the characteristics of worms in the target finding and worm transfer phases into four categories based on the worm's target finding scheme, propagation scheme, transmission scheme, and payload format. Each scheme is further divided into subcategories (Fig. 1). Each one of these categories is discussed in the following sections.

WORM TARGET FINDING SCHEMES

The first step of a worm's life is to find targets. There are many different methods to find the next victim. One simple way is to use blind target scanning, which means the worm has no prior knowledge about the targets. The three types of blind target scanning are sequential, random, and permutation scanning. All these methods are based on chance and have a relatively high failure connection rate. Many worms use this method, and many anomaly-based detection systems are designed to capture this type of worm. Blind scanning worms may be easier to implement and may spread fast, but are not very accurate. The miss rate can be very high. An improved version of the blind scanning scheme is to focus on local subnet scanning with information obtained from the current victim. Doing so can improve the hit rate of scanning.

Although most target scanning worms blindly scan the entire Internet IPv4 address space, an advanced scanning worm, or "routing worm," targets a smaller scanning space without ignoring potential vulnerable hosts. A routing worm uses the information provided by border gateway protocol (BGP) routing tables to narrow the scanning spectrum and target particular systems within a geographic location (e.g., a specific country), an Internet service provider (ISP), or an autonomous system. A routing worm can spread between two to more than three times faster than traditional worms [2].

Furthermore, there is a hypothetical category of scanning worms based on importance scanning.¹ Instead of randomly choosing targets, an importance scanning worm samples targets

¹ Importance scanning is based on importance sampling in statistics, which is used to reduce the sample size for accurately estimating the probability of rare events.

in accordance with the underlying group distribution of the vulnerable hosts. This type of worm usually works in two stages: in the first stage random or routing scanning is used to gather information about enough IP addresses to build an initial group distribution of vulnerable hosts, and then the worm uses an importance sampling technique to reduce the number of scans and attack a large number of vulnerable hosts very fast [3].

With the advent and adoption of Network Address Translation (NAT) and IPv6, researchers have studied their impact on scanning worms. Upgrading to IPv6 can dramatically increase the scanning space (2^{64} IP addresses for a single subnet, compared to 2^{32} IP addresses in the entire IPv4 space!) and, as a result, virtually prevent a worm from spreading through scanning (Zou *et al.* calculated that for a worm with a scan rate of 100,000 hosts/s, it would take 40 years to infect 500,000 vulnerable hosts on a single IPv6 subnet) [2].

In addition, Rajab *et al.* showed that NAT affects and limits worm propagation in three ways. First, it reduces the number of hosts that are globally reachable. Second, if a host is compromised inside the private address space, NAT affects how efficiently the host can discover other vulnerable hosts outside of the private address space. Depending on the scanning technique used by the worm, NAT sometimes limits the worm scan to only the private IP address space. Third, NAT can be a major obstacle for multistage worms where the shell code of the worm on the infected machine needs to download its payload from another victim using a file transfer protocol such as TFTP, while the IP of the other victim will not be globally accessible [4].

The second way of finding targets is to use a prescanned list of vulnerable addresses called a *hit list*. This way, the worm knows exactly where the target is. The hit list can be generated stealthily before the release of a worm or obtained somewhere else. It can be contained inside the worm, or stored somewhere externally for worms to query. The bigger the size of the hit list, the harder it is to obtain and carry, but it is more accurate and may cause more damage. Being pre-generated, a hit list can be out of date since the Internet is changing all the time, but the speed and accuracy of the initial spread will improve greatly over pure blind scan. Because of the high accuracy, the worm will cause very little anomaly on the network and therefore be hard to detect with conventional anomaly-based NIDSs.

Staniford *et al.* [5] simulated a very fast spreading worm named the Warhol Worm that, with a combination of a hit list and permutation scanning, is able to infect most vulnerable systems in possibly less than 15 minutes. The hit list helps the initial spread, and the permutation keeps the speed and hit rate high for longer than just using random scanning. An extended version of the Warhol Worm, which is equipped with a global size hit list, is a flash worm. Staniford *et al.* also simulated this type of worm [6]. The result of the simulation showed that a UDP flash worm can infect 95 percent of one million vulnerable hosts in 510 ms, while a TCP version of flash worm can cause the same damage in 1.3 s.

Many hosts on the Internet store information about other hosts on the network (e.g., /etc/hosts) and possibly reveal their vulnerabilities. Topological worms use this information to gain knowledge of the topology of the network and use that as the path of infection. This makes the attack more accurate since the scanning and infection may look like normal traffic as each infected host needs to contact few other hosts and does not need to scan the entire network. Furthermore, topological worms can spread very fast, especially on networks with highly connected applications [7].

If a worm does not aggressively seek the target but patiently waits, it is said to take the passive approach. Instead of vol-

untarily scanning the network, it waits for potential victims to approach the machine where the worm resides and then replies with a copy of the worm. It may also wait for certain user actions to find the next victim. For example, Gnuman is a passive worm that acts as a Gnutella node waiting for queries to copy itself, and CRClean is a passive worm waiting for an attack from Code Red II to counter it by installing itself on the attacker's machine [7]. This method is slow, but very stealthy and hard to detect.

To avoid being caught by traditional scanning worm detection techniques, a new category of worms has recently emerged that uses popular search engines such as Google and Yahoo to find vulnerable targets. Such worms use carefully crafted queries to find vulnerable hosts on the Internet. As an example, a new worm called "Santy" uses Google to search for Web servers that contain the string "viewtopic.php" to exploit a vulnerability in phpBB2 and infect the Web servers [8].

WORM PROPAGATION SCHEMES

After the next victim is found, a copy of the worm is sent to the target. There are different schemes for worm propagation. In [7] three propagation schemes are mentioned: self-carried, second channel, and embedded.

In self-carried worms, propagation is straightforward; the worm payload is transferred in a packet by itself. Other worms are delivered through a second channel; that is, after finding the target, the worm first goes into the target, and then downloads the worm's payload from the Internet or a previously infected machine through a backdoor, which has been installed using RPC or other applications. A more deceitful worm may append the payload after, or replace, legitimate traffic to hide itself. This embedded propagation scheme is very stealthy. No anomalous events will be triggered, and it is hard for anomaly-based detection systems to detect. Contagion strategy is an example of a worm that uses embedded propagation [7].

In addition to the three propagation schemes discussed, botnets have been utilized to propagate worms, spams, spyware, and launching distributed denial-of-service (DDoS) attacks [9]. A botnet is a group of compromised hosts under the control of a botmaster. The communication channel for the botmaster to issue commands can be implemented using different protocols such as http or point-to-point (P2P) protocols; however, the majority of botnets use the Internet Relay Chat (IRC) protocol for this purpose [10]. Witty is an example worm propagated by botnets. Witty infected 12,000 vulnerable hosts in 45 min, and when the machine that launched the initial attack was discovered, it was not subject to the vulnerability exploited by the Witty worm, instead, a different vulnerability was used to take over the control of the machine by a botmaster to launch the attack [9].

Full treatment and defense against botnets themselves are outside the scope of this article. However, some of the methods we discuss later, such as combining the intelligence of the control and data planes by Zhang *et al.* [11], can also be used to combat botnets.

WORM TRANSMISSION SCHEMES

Based on how worms are transmitted, there are TCP worms and UDP worms. The major difference between these two types of worms is that TCP worms are latency-limited and UDP worms are bandwidth-limited.

All TCP connections require a three-way handshake to establish connection before transmission. Therefore, after a host sends out a TCP SYN packet to initiate a connection, it

must wait until it receives a corresponding SYN/ACK or timeout packet from the other end before it can take any further actions. Compared to UDP worms, TCP worms need an additional round-trip time and the two 40-byte packets to establish the connection. During this wait time, the thread or process is blocked and cannot infect other hosts.

UDP is connectionless, so UDP worms do not require a connection to be established before infection can begin. The implementation of the worm is normally self-carried and included in the first packet sent to the target. Since there is no wait time required as for TCP worms, UDP worms normally spread very rapidly, and their speed is only limited by network bandwidth. UDP worms often have to compete with each other for network resources [12].

WORM PAYLOAD FORMATS

The term *payload* used here means the actual worm code. Traditionally, worms send the payload in a straightforward unchanged fashion. By matching the worm payload with the signatures in a database, signature-based detection systems can identify them. Some worms make the payload variable size by padding the payload with garbage data, but the signature will not change; this is still a *monomorphic* worm.

Worm authors can make changes in the payload to make them appear innocent to evade detection systems. They may fragment worm payload differently and reassemble the pieces at the target. This type of worm is also classified as *monomorphic*. The term *polymorphic* used here describes those worms that change their payload dynamically by scrambling the program, so every instance of the worm looks different but functions exactly the same way. With changing appearances of worms, it is very hard for traditional signature-based detection systems to detect them.

If a worm can change not only its appearance but also its behavior, it is a *metamorphic* worm. If the worm also uses a complicated encryption scheme to hide its true purpose, it will be even harder to defend against [13].

EXISTING INTERNET WORMS

In this section we look at one of the first Internet worms, the Morris worm, which gained extensive media coverage, then discuss five more recent Internet worms: Code Red, Nimda, Sasser, Slammer, and Witty, based on their characteristics.

The Morris Worm — The Morris worm was one of the first Internet worms whose devastating effect gained the wide attention of the media. Morris worm was launched in November 1988 by Robert Tappan Morris, who was a student at Cornell University at the time. It is the first known worm to exploit the buffer overflow vulnerability. It targeted sendmail and finger services on DEC VAX and Sun 3 hosts [1].

Based on the creator's claim, the Morris worm was not intended to cause any harm, but was designed to discover the number of the hosts on the Internet. The worm was supposed to run a process on each infected host to respond to a query if the host was infected by the Morris worm or not. If the answer was yes, the infected host should have been skipped; otherwise, the worm would copy itself to the host. However, a flaw in the program caused the code to copy itself multiple times to already infected machines, each time running a new process, slowing down the infected hosts to the point that they became unusable [14].

Code Red I and Code Red II — Code Red I was first seen in July 2001 affecting computers running Microsoft's Internet

Information Server (IIS) Web service. In the first 20–25 days after getting into the machine, Code Red I uses a blind scan scheme that scans port 80 on random IP addresses to find other vulnerable machines, and then it launches a denial-of-service (DoS) attack targeting a set of IP addresses. The infected websites will display: "HELLO! Welcome to <http://www.worm.com>! Hacked By Chinese!"

Code Red II was released one month later. It is a variant of the original Code Red. Code Red II no longer launches a DoS attack against predefined IP addresses; instead, it installs a backdoor into the infected systems. It still employs blind scan but focuses more on the local subnet, and targets mainly systems with Chinese language settings.

Code Red I sends its payload in monomorphic format and has a signature starting with "GET /default.ida?NNNNNNN." Code Red II has a similar signature, but replaces N with X. Both versions of Code Red are self-carried and transfer via TCP connections.

Nimda — Nimda was first reported in September 2001, targeting Microsoft Windows workstations as well as servers.

Nimda is an advanced multivector worm, which uses multiple mechanisms to spread itself, including from client to client via email, from client to client via network shares, from Web server to client via browsing compromised Web sites, from client to Web server by active scanning for various vulnerabilities of Microsoft IIS 4.0 / 5.0, and from client to Web server by scanning for the back doors installed by Code Red II and Sadmin/IIS [15, 16].

Nimda propagates itself by sending emails to anything that looks like an email address inside .htm or .html files in a user's Web cache folder as well as contents of a user's email messages retrieved via the MAPI service. The subject line of the message is variable, and the attached binary also has some variations resulting in different MD5 checksums, but all binaries are exactly the same size (57,344 bytes). In addition, Nimda scans the network to find and infect vulnerable IIS servers on TCP port 80 as well as using UDP port 69 to download the worm to IIS via TFTP [16].

Once a host is infected, Nimda allows the attacker to run commands with the same privileges of the infected user as well as using the infected host as a zombie to participate in DoS attacks on third parties. In addition, the high scanning rate of Nimda can result in bandwidth DoS attacks on networks with infected hosts.

Furthermore, Nimda replaces existing binaries on the system with Trojan horse copies and infects all Web contents such as .htm, .html, and .asp on the system, so any user browsing these contents on the system via a Web browser will download a copy of the worm; in some cases, certain browsers will execute the code automatically, infecting the user's host [16].

Slammer/Sapphire — Slammer, also known as Sapphire, was one of the smallest worms ever seen. It was found in January 2003 targeting Microsoft SQL Server 2000 or MSDE 2000.

Slammer uses UDP port 1434 to exploit a buffer overflow in an MS SQL server. The code size is 376 bytes. Adding the UDP header makes the worm 404 bytes long in total [12]. It uses a blind scan scheme where randomly generated numbers are used as IP addresses in searching for vulnerable hosts. To initialize the random number generator, Slammer uses the GetTickCount() function from Win32 API. Sometimes the random generator returns values that are broadcast addresses, such as a.b.c.255, and causes all the hosts in that network to receive the worm packets, making the spread of the Slammer worm more rapid. Like most UDP worms, Slammer is self-

	Target finding scheme	Propagation scheme	Transmission scheme	Payload format
Morris	Blind	Self-carried	TCP	Monomorphic
Code Red	Blind*	Self-carried	TCP	Monomorphic
Nimda	Blind	Self-carried	TCP and UDP	Monomorphic
Slammer	Blind	Self-carried	UDP	Monomorphic
Sasser	Blind	Second channel	TCP	Monomorphic
Witty	Blind	Botnet	UDP	Monomorphic
*Code Red II focuses on local subnet scan				

■ Table 1. Existing Internet worm implementations.

carried and has a monomorphic payload.

Slammer does not write to the disks of infected machines, it only overloads the victim systems and slows down traffic [17].

Sasser — Sasser was released in April 2004 targeting systems running Microsoft Windows XP or Windows 2000 that had not been patched for the vulnerability of Local Security Authority Subsystem Service (LSASS). Sasser exploits a buffer overflow vulnerability of LSASS to gain access to remote systems and spread further. Sasser transfers with a second channel via TCP connection and uses a monomorphic payload.

If Sasser successfully infects a system, it will act as an FTP server listening on TCP port 5554. Sasser then generates 128 scanning threads (Sasser B uses processes instead of threads) to find vulnerable systems using random IP addresses. The worm probes and tries to connect to the next victims through TCP port 445, then attempts to connect to the victim's command shell available on TCP port 9996. Once the connection is successful, the victim will download the worm code from the attacker using FTP. The sizes of Saspers A through E are 15–16 kbytes. Sasser F and later versions are larger; Sasser F is 74 kbytes [18], and Sasser G is 58 kbytes [19].

After the Sasser worm enters a system, it makes a copy of itself, stores one copy in the Windows directory, and adds itself to the Registry. Transactions through the FTP server are logged to C:\win.log.

The Witty Worm — The Witty worm was released in March 2004, targeting buffer overflow vulnerability in several Internet Security Systems (ISSs), including RealSecure Server Sensor, RealSecure Desktop, and BlackICE. Witty took advantage of a vulnerability of the ISS Protocol Analysis Module (PAM) used for ICQ instant messaging.

Witty is a self-carried monomorphic UDP worm that employs a blind target finding scheme. It sends out UDP packets to 20,000 random generated IP addresses on random destination ports from source port 4000, with a random packet size ranging between 768–1307 bytes. The code size of Witty is only 637 bytes, and the rest of the payload is padded with data from system memory. This padding does not change the monomorphic format of Witty. The payload contains the text “(^.^) insert witty message here (^.^),” which is why it is named Witty. Witty randomly writes data onto the disk of infected machines [20].

It is harder to detect Witty worms than worms with fixed size packets targeting fixed destination port numbers because of its random characteristics. The size of Witty worms is larger than Slammer worms (some can be doubled), but they

spread faster than Slammer. This proves that size is not always the bottleneck for the spreading of UDP worms [21].

Another significance of the Witty worm, as mentioned earlier, is that Witty was one of the first known worms distributed using botnets [9].

We summarize the characteristics of the above worms in Table 1.

MULTIPLATFORM AND MULTIEXPLOIT WORMS

Up until now, most worms have attacked only one type of operating system, and often target a single vulnerability, so network administrators only need to patch one type of system after receiving the worm alert. If a worm can perform multiplatform attacks, it will be harder and more complicated to defend against. And since there is more work to be done, the response time will be slower as well.

If a worm infects victims by using multiple vulnerabilities, a single worm can cause greater damage at a faster rate, and the work of patching will also be harder as well. This presents a challenge in worm research.

DEFENDING AGAINST INTERNET WORMS

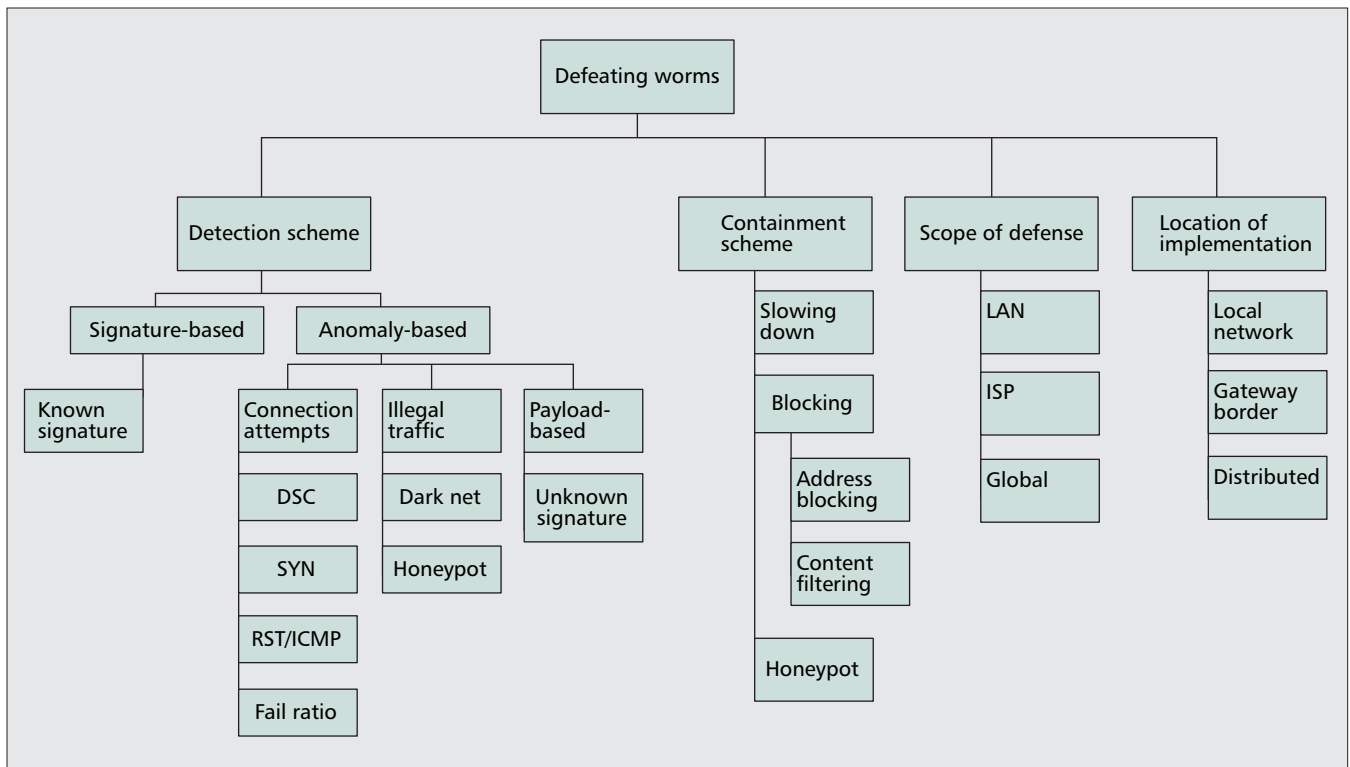
Now that the nature of Internet worms is known, the next question is how to defend against them. This task breaks down into worm detection and worm containment. Detection and containment systems can be implemented at different locations in the network, and thus have different scopes of defense. We categorize worm defense schemes in Fig. 2 and discuss each subcategory in the remaining sections of this article.

WORM DETECTION

Based on the parameters used for detection, detection algorithms can be roughly divided into signature-based and anomaly-based schemes, as seen in Fig. 2. There are many proposed algorithms for both schemes. This section first introduces signature-based detection and then discusses anomaly-based detection.

SIGNATURE-BASED

Signature-based detection is a traditional technique used for intrusion detection systems (IDSs) and is normally used for detecting known attacks. There are different definitions of



■ **Figure 2.** *Categorization of Internet worm defense.*

worm signature. In this article our discussion will focus on the content signature, which is often a string of characters that appear in the payload of worm packets as part of the attack.

No knowledge of normal traffic is required, but a signature database is needed for this type of detection system. This type of system does not care how a worm finds the target, how it propagates itself, or what transmission scheme it uses. Signature-based systems take a look at the payload and identify whether or not it contains a worm. Since every packet is examined, signature-based systems can catch worms that employ self-carried or second channel propagation schemes. Embedded worms may not be detected because the payload can be different from worm to worm, depending on the embedding method used.

One big challenge of the signature-based IDS is that every signature requires an entry in the database, so a complete database might contain hundreds or even thousands of entries. Each packet is compared to all entries in the database. This can be very resource-consuming, and doing so will slow down the throughput, making the IDS vulnerable to DoS attacks. Some IDS evasion tools use this vulnerability and flood signature-based IDSs with too many packets to the point that the IDS cannot keep up with traffic, thus making the IDS time out and miss packets, and, as a result, possibly miss attacks [22]. Furthermore, this type of IDS is still vulnerable against unknown attacks.

We believe the deficiencies of the signature-based detection method can be addressed by incorporating an anomaly-based unknown signature detection scheme with signature-based detection in a two-tier architecture, supported by an aging and removal process to keep the size of the signature database small. The signature-based detection engine can run on a small and efficient database to look for any known threats, and at the same time the anomaly-based unknown signature detection scheme can work slower on the traffic and provide signatures for any new threat to the IDS's database. The aging process ensures removal of old signatures not seen for a long time from the database to keep the database small

and the process as efficient as possible; if an old worm resurfaces again, it will be detected by the unknown signature detection engine and added back to the database.


ANOMALY-BASED

The signature of a new worm is unknown before it is seen, and it is difficult to draw conclusions based on a small number of packets. All fast spreading worms seen so far create large traffic volume, and most of them employ blind scan. Lots of scans target nonexistent addresses and closed ports. Since most networks behave in a particular and consistent fashion most of the time, when there are lots of abnormal phenomena occurring on the network, it often means something is wrong.

Signature-based detections check the payload of worms to generate and match signatures. Most anomaly-based detections do not care about the payload format or content; instead, they check the headers of packets to define the type of connection to which the packet belongs. They observe the network traffic volume and the monitored hosts' behavior. The three most common purposes for a packet sent or received by a host are to initialize a connection, indicate a failed connection attempt, and send data via an already established connection. The system may also keep track of the traffic between source and/or destination addresses and try to find the scanner.

While a packet's header information is useful to detect attacks exploiting vulnerabilities of the network stack or probing hosts for vulnerable services, a packet's payload information can be used to detect attacks directed at vulnerable applications. There is a category of anomaly-based detections that examines the packet payload to detect attacks directed at applications as the connection in these types of attacks is normally established (so checking the headers would not reveal the attack). An example of this anomaly-based NIDS is POSEIDON [23].

In general, anomaly-based detection systems detect abnormal behaviors and generate alarms. This technique often



■ **Figure 3.** Connection attempts: a) successful TCP connection; b) TCP destination port closed; c) UDP destination port closed; d) destination IP address does not exist.

requires the definition of normal network behavior, which depends on a training period before the system can be effective in protecting the network. If an attack is crafted carefully, it may possibly train the system to take an anomaly as normal or trigger false alarms. While this method is generally the best approach to detect unknown worms, the big challenges of anomaly-based detection systems are defining what normal network behavior is, deciding the threshold to trigger the alarm, and preventing false alarms. The users of the network are normally human, and people are hard to predict. If the normal model is not defined carefully, there will be lots of false alarms, and the detection system will degrade performance.

As seen in Fig. 2, anomaly-based algorithms are categorized based on connection attempts, illegal traffic, or packet payload. Using connection attempts, the schemes may rely on connection count/traffic rate, failure connection rate, success/failure ratio, or destination-source correlation in their detection. Using illegal traffic, schemes may monitor darknet and honeypots. Using a packet's payload, the schemes try to measure anomaly by comparing the packet payloads to a reference model made during normal traffic (training period). The following subsections explain these categories in detail.

Traffic Rate/Connection Count: TCP SYN — Worms send out large numbers of scans to find victims. Keeping track of the outbound connection attempts is a traditional way to detect scanning worms. TCP/IP protocol stacks require the host to send out a TCP SYN packet to initiate a connection (Fig. 3a), and this is used as the parameter for connection count detection. The idea is that if the number of SYN packets sent from a certain host exceeds a threshold value within a period of time, the host is considered to be scanning.

This method is used in many older algorithms and some commercial IDSs, such as the older version of snort [24]. Tracking TCP SYN packets may be able to catch most active scanning worms with most scanning schemes, but it is very easy to cause false alarms. It is not widely used nowadays because it is not very accurate, nor is it efficient. Also, if the system logs TCP SYN only, it is useless against UDP worms. Due to inefficiencies and high numbers of false alarms, this


method is for the most part obsolete, and we do not recommend using it.

Connection Failure Rate: TCP RST and ICMP — Many worms implement blind target finding schemes. No matter whether the worms use sequential scan, permutation techniques, or randomly generated IP address and/or random ports scan, the scan will include unused address spaces and closed ports, causing the connection attempt to fail. While ordinary users access the Internet mainly via domain names, they are less likely to encounter failed connection attempts.

Attempting to connect to a nonexistent IP address or an existing address with the target port closed, the connection attempt is considered failed. According to TCP/IP, if the destination host does not exist, an ICMP host unreachable packet is returned (Fig. 3d). A TCP RST packet is returned when a TCP connection targets an existing host with the destination port closed (Fig. 3b), and an ICMP destination unreachable packet error message is returned if this is a UDP connection (Fig. 3c). With these characteristics of TCP/IP, keeping track of these error messages will work well against blind target finding schemes. The scan is blind, so the failure rate is higher than normal. In the case of hit list, topological, or passive scanning worms, failure rate will not be very useful because the scan and spread of worms have valid targets and will not cause these error messages.

Compared to the traditional method of detecting TCP SYN packets for connection count, it is more efficient and accurate to detect active scanning worms depending on failed connections. This approach may be useful for both TCP and UDP worms. ICMP error messages will be sent for both TCP and UDP connection attempts, but TCP RST works with TCP connections only. However, this scheme becomes less effective if the ICMP error messages are blocked or dropped by some border routers or gateway systems.

Berk, Bakos, and Morris [25] proposed a global detection algorithm based on ICMP destination unreachable error messages. Routers often generate ICMP error messages to notify the source that the target IP address does not exist on the network. By forwarding these messages to a central collection point, an alert can be generated when the number of such



■ **Figure 4.** Illustration of a destination-source correlation scheme.

error messages reaches a certain threshold. The Distributed Anti-Worm (DAW) architecture [26] identifies scanning sources by keeping track of both TCP SYN and RST packets, dealing with TCP worms only.

If the source address is forged in the packet header, it will be very difficult to detect the scan source. This can be used as a feint attack by worm authors. If the system administrator actually traces back the fake information given in the header, it will be a waste of effort. The worm may also use this technique to trigger false alarms. Forging header information is used in many Internet attacks. Although this technique is not commonly seen in worm attacks yet, it is still an issue worth attention. The Worm Early Warning System (WEW) proposed by Chan and Ranka [27] considered this problem. The proposed architecture utilizes gateways and hash algorithms to not only detect the error messages from failure connection attempts, but also verify whether the source address is legitimate or forged. The system only takes actions on failed connection attempts that are sent from existing source addresses.

Another issue of detection schemes based on the connection failure rate is that the worm might initiate thousands of connections before enough failures are observed. Schechter *et al.* proposed an algorithm based on a combination of the Reverse Sequential Hypothesis Testing algorithm to monitor the connection failures and the Credit-Based Connection Rate Limiting (CBCRL) algorithm to limit the rate, in which the first contacts can be initiated by each host on the local network [28].

Furthermore, any method relying on connection status requires resources to keep track of hosts and connection information, which means such a method is not suitable for large networks.

Ratio of Success and Failure Connections — Instead of counting the failure or successful connection attempts, some believe it is the ratio or correlation of successful and failed connections that matters. Counting the number of connections, whether successful or not, depends on the Internet usage and network size to be effective. If usage is too low or the network too small, using connection counts as a detection parameter may be less accurate. Thus, both success and failure connections should be taken into account. When the percentage of failed connections is large enough, this is said to be anomalous, and an alert will be generated. Similar to the pre-

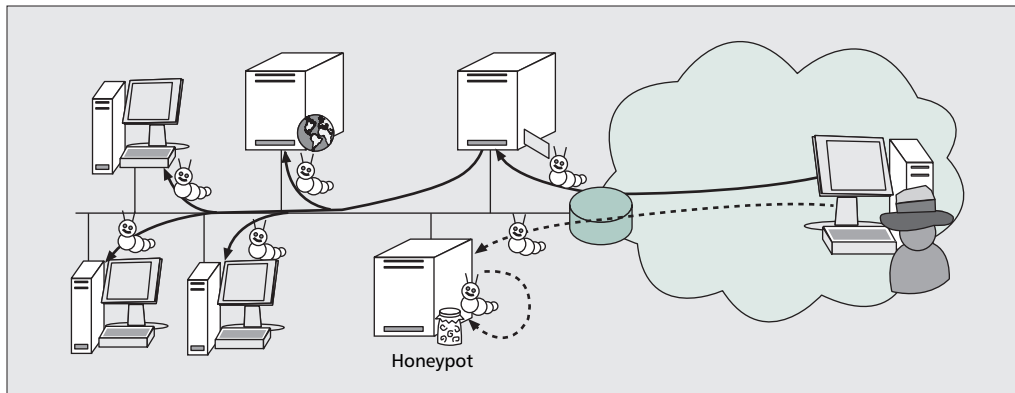
vious method of detecting failure connection attempts, this method works well against blind scans, but not with other scanning techniques where the targets are specific and legitimate.

Jung *et al.* [29] proposed the Threshold Random Walk (TRW) algorithm, derived from Sequential Hypothesis Testing. The algorithm says that for a given remote host R , if it tries to make a connection to a local host L , this attempt can be a success (marked 0) or a failure (marked 1). With a sequence of these results of connection attempts, the system can decide whether the remote host is a scanner based on the test of hypothesis. This algorithm requires very few packets (only four or five) to draw a conclusion and does not require training of the system in advance. It focuses on detecting TCP traffic only.

Weaver *et al.* [30] introduced the concept and importance of hard LANs. They explained the three limitations of the TRW algorithm: it is offline, requires infinite states, and requires potentially unlimited memory access time. Because of these limitations, they proposed an improved algorithm that, instead of identifying a connection as completed or failed, considers all new connections as failures and changes the status to success when there is a response. It also keeps track of UDP connections. This “guilty until proven innocent” method keeps a counter for every IP address, starting with a miss. If the connection succeeds, the corresponding counter decrements. The counter increments if the connection is failed. When a counter is greater than 10 or another predefined value, the corresponding system is considered a scanner. This algorithm requires a fairly small amount of memory and is suitable for integration into switches or other low-cost networking devices.

Monitoring the connection status, whether successful, failed, or both, often requires keeping track of each host’s and/or each connection’s information. If the network being monitored is large, this can be very resource consuming.

Destination-Source Correlation — The Destination-Source Correlation (DSC) algorithm is a two-phase local worm detection algorithm that aims to detect fast spreading scanning worms [31, 32]. Instead of watching for connections and the ratio of successful and failed attempts, this algorithm is based on the correlation between incoming and outgoing traffic. DSC keeps track of SYN packets and UDP traffic of the



■ **Figure 5.** Honeypot used in worm detection and containment.

source and destination. It is illustrated in Fig. 4, where for every port, if a host inside the monitored network previously receives a packet on a certain port (e.g., port 25 in the illustration) and then starts sending packets designated to the same port on which it previously received packets, a counter is incremented. When the counter reaches a certain threshold, an alert is issued.

The DSC algorithm is able to detect almost all types of scans, as long as the scan is frequent enough (based on the threshold) and the infection of the worm is targeting the same port. It can detect aggressive scans including blind and topological scans. The effectiveness for passive scan depends on the incoming traffic rate since it relies on interaction with the worm. It works for both TCP and UDP worms.

The major issue of DSC is that it can only capture scans from worms targeting the same port. To address this issue, Qin *et al.* [32] combined HoneyStat with a modified version of DSC. Based on the DSC algorithm, the system monitors IP or medium access control (MAC) addresses to defend against worms using IP spoofing. HoneyStat is used to gather statistical data about the attack. Since DSC only captures scans with the same port and HoneyStat can capture scans with different ports, HoneyStat can cover what DSC cannot see.

In addition, the EarlyBird system combines an unknown signature-based algorithm with a DSC-like algorithm for worm detection [33]. An alarm is generated when packets with similar contents are sent to a number of destination IP addresses, are received from a large number of source IP addresses, or are sent from a number of hosts to a large number of hosts (source and destination IP address pairs). Stanford *et al.* [6] also indicated that this type of system may be able to capture Flash worms.

DarkNet/Unused Address Space — Worms using blind scan generate random numbers for target addresses. There is a very high chance that these addresses are unused. Monitoring unused address space instead of used ones is another approach. This is a branch of anomaly-based detection since scanning or connection attempts toward nonexistent addresses are abnormal behaviors of a regular network.

The monitored address space has to be big enough for this method to be useful. Chen, Gao, and Kwiat [34] presented the Analytical Active Worm Propagation (AAWP) model, which simulates the propagation of random scanning worms. Using this model, they derived the size of address space needed to detect active worms. They suggested that an address space of 2^{24} IP addresses is large enough to detect worms effectively, and an address space smaller than 2^{20} addresses will be too small to obtain a realistic result of the spread of worms.

A scanner host is normally a host infected by worms. In other words, it is a victim itself. Wu *et al.* [35] proposed an algorithm based on the number of victims. The victim is

defined as: “the addresses from which a packet is sent to an inactive address” [35]. This means if an IP address send packet to an unused IP address, this source IP address is considered a victim. To prevent false alarms, they combined this definition with the Two Scan Decision Rule (TSDR), which means if the system captured two packets sending to unused IP addresses from the same host, the host is a victim. When victim count reaches a certain threshold, the system will generate a worm alert.

One of the major advantages of this method is that it requires significantly less resources than methods looking at the normal traffic in the used address space (as there should be normally no traffic in an unused address space), and records much less information in the IDS database.

Monitoring an unused IP address can find worms using blind target finding schemes. But again, it is not very useful against hit list, topological, or passive scans. This method works for both TCP and UDP worms since both transmission schemes require IP addresses.

Honeypots — Honeypot technology can be used for anomaly-based worm detection. A honeypot is a vulnerable system on the network that does not provide any real services. In [36] Spitzner defines the honeypot as a “security resource whose value lies in being probed, attacked, or compromised.” Figure 5 illustrates the setup of a honeypot where it appears as a normal vulnerable machine on the same network, just like other servers and hosts, to lure attackers.

There are many different implementations based on the level of interaction the honeypot provides. In a normal situation, no traffic is supposed to come toward the honeypot; therefore, any traffic targeting the honeypot is considered anomalous and may be an attack. Compared to other IDSs, honeypot systems gather less but higher-quality data because every piece of data is information on probing or attack.

Honeypot can detect blind scan worms for the same reason as the approach of monitoring unused address spaces. Honeypot can also defend against hit list scanning worms. If the hit list is generated automatically, the prescan may very likely include a honeypot system because it appears as a normal vulnerable host on the Internet. Honeypot can be useful against topological worms if other working hosts on the network contain proper information on the honeypot and let the worm find it. But honeypots are not useful against passive worms because they only sit and wait, but do not initiate connections. As long as they are properly configured, honeypots can detect both TCP and UDP worms.

Virtual Honeypot was used for worm detection [32, 36, 37]. In an emulator they created a minimal honeypot that uses virtual machines and multihome to cover a large address space, called HoneyStat. It is used to gather information about worms as well as capture worms. Their HoneyStat simulation

runs on VMware GSX, so if there are 64 virtual machines running Windows, and each has 32 IP addresses, a single node can have 211 IP addresses. The hardware requirement for such a system can be as low as 32 Mbytes RAM and 770 Mbyte virtual drives to capture worms. HoneyStat generates alerts based on the correlation of three types of events: memory, disk write, and network events. An example of a memory event is buffer overflow. An example of a network event is downloading of some malicious code. An example of a disk event is writing to registry keys or critical files.

HoneyD [38], a low interaction open source HoneyPot daemon that supports both UNIX and Windows platforms, can detect and log connection on any TCP or UDP ports. When a connection to HoneyD is established, HoneyD will emulate the configured personality or operating system and port behavior based on the configuration script. HoneyD can emulate any of the 437 existing operating systems and any size of network address space with desired topology. Provos [39], the author of HoneyD, took HoneyD and built a system on top of it for worm detection. It can detect intrusion as well as be configured to replay incoming packets to higher interaction honeypots for analysis of unusual activities.

A honeypot that uses scripts is more flexible than one with limited configuration settings. But no matter what, a honeypot has a narrower view, since it can only see the traffic coming toward the addresses it simulates. Honeypots can also be used for containment, as discussed later.

Honeypots are most useful when combined with other IDS methods. For example, the Honeycomb project complements a honeypot-based IDS system with generation of signatures to detect unknown worms. Honeycomb uses the Longest Common Substring (LCS) algorithm to detect similarities and patterns in the packet payloads of the traffic seen on the honeypot [40].

Anomaly Detection Systems Based on Packet Payload —

The anomaly-based detection systems discussed so far do not use packet payload information. While a packet's header information is useful to detect attacks exploiting vulnerabilities of the network stack or probing hosts for vulnerable services, a packet's payload information can be used to detect attacks directed at vulnerable applications since the connection in these types of attacks is normally established, and checking the headers would not reveal the attack [23].

Packet Header Anomaly Detection (PHAD) is a partly payload-based system, which learns the normal ranges of values for each packet header field at the data link (Ethernet), network (IP), and transport/control layers (TCP, UDP, ICMP); however, PHAD does not check application layer protocols. PHAD examines 33 packet header fields, and is designed to be as protocol independent as possible. PHAD, like all anomaly-based systems, checks for unusual events and uses ranking systems to decide how unusual they are. The rarer they are, the more likely they are to be hostile. PHAD uses the rate of anomalies during the training period to estimate the probability of anomalies in network traffic. Based on this information, PHAD calculates a score for each packet header inversely proportional to the probability of being anomalous. At the end, the scores for all 33 packet header fields are added up to calculate the final score of the packet to help decide if the packet should be considered anomalous [41].

PAYL is a fully payload-based system based on modeling normal payloads expected to be delivered to the network specific to the site at which PAYL is installed. During the training period, PAYL creates a profile based on the traffic to each service during normal operation, and produces a byte

frequency distribution as a model for normal payloads. Based on this information, a centroid model is created prior to the anomaly detection phase for each service. In the anomaly detection phase, the distance of each packet payload from the centroid model is calculated, and if the payload is too distant from the normal payload, it will be considered anomalous. The main difference of PAYL in comparison to PHAD is that PAYL looks at the whole payload rather than the 33 packet header fields, so it can also detect application level anomalies. Furthermore, PAYL clusters the centroids to increase accuracy and dramatically reduce resource consumption [42].

POSEIDON is one of the most recent anomaly detection systems based on packet payload. POSEIDON has a two-tier architecture, including a self-organizing map (SOM) as a pre-processor to classify the payload data and a slightly modified PAYL system. A SOM is a topology-preserving single-layer map, which preserves the neighborhood relation between nodes during classification. It requires some parameters on startup including total number of nodes on the network. The modification in the PAYL system used in POSEIDON is the preprocessing of packets by the SOM. Damiano Bolzoni *et al.* [23] report that POSEIDON has a higher detection rate, a lower number of false positives, and higher runtime efficiency than PAYL and PHAD [23].

Detecting Polymorphic Worms Based on Unknown Signature Detection Systems —

We mentioned previously that one of the limitations of signature-based detection systems is vulnerability against unknown attacks. To remedy this issue, some algorithms have been proposed to detect unknown attacks by generating signatures in real time. These algorithms are considered anomaly-based as they generate signatures based on what they detect to be a worm when analyzing the network traffic rather than using existing signatures to detect worms. As discussed in the previous section, PAYL is a good example of such a system, capable of automatically generating signatures for unknown worms [43].

Madhusudan and Lockwood [44] introduced an algorithm to detect frequently occurring strings in packets and use them as signatures to use for detection. In this system a signature detection device (DET) sits between the router and subnets, and monitors the traffic flow to detect Internet worms. It is implemented in hardware, and throughput is improved by parallelism and hashing. The EarlyBird system used a similar approach to find frequently occurring substrings in packets [33]. As the algorithms detect substrings in a worm signature, they can not only detect unknown worms sent in one packet, but also be effective against worms sent in fragments across several packets, even when worms break the payload differently each time. In addition, it can catch worms that are embedded in legitimate packets for propagation.

We also discussed honeypots and how they can be used as non-payload anomaly-based detection systems. Furthermore, honeypots can be also used to generate signatures for unknown worms. Honeycomb is a honeypot-based IDS system that is capable of generating signatures for unknown worms. Honeycomb deploys the Longest Common Substring (LCS) algorithm to spot similarities and patterns in packet payloads of the traffic seen on the honeypot [40].

No matter whether the signature is known or unknown, most detection algorithms target monomorphic worm payloads only and have no defense against polymorphic worms, which change the payload dynamically. Kim and Karp [45] proposed *autograph*, a distributed worm signature detection system capable of dealing with polymorphic and potentially metamorphic worms. Autograph relies on unsuccessful scans to identify suspicious source IP addresses and segregates flows

Characteristic of worms Method of detection		Target finding scheme				Propagation scheme			Transmission scheme		Payload format		
		Blind	Hit list	Topo-logical	Passive	Self-carried	Second channel	Embedded	TCP	UDP	Monom-orphic	Poly-morphic	Meta-morphic
Signature	Known signature	—	—	—	—	✓	✓	Maybe	✓	✓	✓	x	x
	Token-based signature	—	—	—	—	✓	✓	✓	✓	✓	✓	✓	x
Anomaly-based	Destination-source correlation	✓	✓	✓	x	✓	—	—	✓	✓	—	—	—
	TCP SYN — connection count	✓	✓	✓	x	✓	—	—	✓	Maybe	—	—	—
	Failed connection attempts	✓	x	x	x	✓	—	—	✓	Maybe	—	—	—
	Ratio of success/failure attempts	✓	x	x	x	✓	—	—	✓	Maybe	—	—	—
	Monitoring dark net	✓	x	x	x	✓	—	—	✓	✓	—	—	—
	Honeypot	✓	✓	Maybe	x	✓	—	—	✓	✓	—	—	—
	Payload-based (unknown signature)	—	—	—	—	✓	✓	Maybe	✓	✓	✓	✓	Maybe
Hybrid	Modified DSC + HoneyStat [21]	✓	✓	✓	x	✓	—	—	✓	✓	—	—	—
	Early Bird (unknown sig+DSC)	✓	✓	✓	x	✓	✓	✓	✓	✓	✓	x	x

■ Table 2. Anomaly detection methods vs. worms characteristic.

by destination port. It automatically generates signatures for TCP worms by analyzing the contents of the payload based on the most frequently occurring byte sequence in the suspicious flow. Autograph consists of three modules: a flow classifier, a payload-based signature generator, and a tattler. A tattler is a protocol based on RTP Control Protocol (RTCP), which facilitates sharing suspicious source addresses among all monitors distributed across the network [45].

Autograph still relies on a single contiguous substring of a worm's payload of sufficient length to match the worm, and the assumption is that this single payload substring will remain invariant on every worm connection; however, a worm in theory can substantially change its payload by encoding and re-encoding itself on each connection to evade being detected by a single substring [46]. To address this problem, Newsome, Karp, and Song [46] proposed *polygraph*, an algorithm to automatically generate signatures for polymorphic worms without a single payload substring. They found that even though polymorphic worms change the payload dynamically, certain contents are not changed. Such contents include protocol framing bytes (e.g., GET and HTTP protocol indicator) and the value used for return address or a pointer to overwrite a jump target. Based on this characteristic of polymorphic worms, they divide signatures into tokens. The system generates tokens automatically and detects worms based on these tokens. An algorithm that detects polymorphic worms can detect monomorphic worms as well, but not the other way around, so it is a more thorough approach.

Unknown signature detection generates signatures from the traffic flow. Even though it takes time to generate signatures, compared to known signature-based detection systems, it may be less efficient when facing known worms, but it can detect newly released worms and possibly catch other kinds of Internet attacks such as DDoS attacks. It may also help in detecting embedded worms since the signature can be part of the packet payload instead of the whole content. Unknown

signature detection systems often do not store all the signatures. As a previously generated signature ages, it will eventually be eliminated from the database, so the database does not just grow bigger and bigger. This conserves system resources from the processing time of comparing signatures and the storage space of the database.

Detecting Search Worms — The techniques used to detect scanning worms (e.g., TCP/SYN, connection failure/success rate) do not work for search worms. Also, signature-based systems are not suited to detect search worms as different queries can produce the same result. Provos *et al.* proposed a solution based on the polygraph framework, which is not dependent on search queries but instead looks for search results. If it finds a particular query returns too many vulnerable hosts (which are tagged during indexing), it removes the vulnerable results from the return list, hence stopping the spread of the worm [8].

LIMITATIONS, BENEFITS, AND COMBINATION USAGE OF DETECTION SCHEMES

So far in this section we have discussed various algorithms in worm detection. These systems can be classified as signature-based or anomaly-based, and are further organized into several subcategories based on algorithms. Different detection schemes are useful against different worm characteristics. This is summarized in Table 2.

An unknown signature-based detection system may defend against zero-day attacks as well as known worms. But it takes time to generate signatures, and since there are defined signatures already, why not just use them? A system equipped with a known worm signature database and an additional real-time signature generator may be more comprehensive and efficient, and it would be even better if it also has the capability of detecting polymorphic worms.

A worm can be detected on the network during the phase of target finding and transmission. Different detection methods catch different types of worms, and no single current algorithm is perfect. A hybrid system with integration of both anomaly-based and signature-based types of detection techniques will give a broader and more complete view to a detection system. Ideally, a hybrid system should check for both signature and network anomalies, and have a honeypot to aid detection and gathering of worm information.

As shown in Table 2, no single algorithm provides complete protection against worms with different characteristics. Most anomaly-based systems focus on detecting blind scan worms, which is by far the most commonly seen technique for active worms. There is no algorithm that can detect passive scanning worms because these worms do not trigger any error messages and normally do not cause high traffic volume. Worms that use embedded propagation schemes are harder to detect because if the worm payload is appended to the packet content, the signature, as a whole, may be different for each worm, unless the system breaks the signatures into pieces and inspects them individually with consideration of the other pieces. Passive scan and embedded payload are often used together in worm implementations. Doing so makes the worm very stealthy, but its spread is slow and results in less damage, making it less of a threat.

There are not many systems that are able to catch worms using topological or hit list scanning schemes because they cause few or no failure connections. However, it is not impossible since they often still generate large traffic volume. Antonatos *et al.* proposed Network Address Space Randomization (NASR), a solution based on the concept of frequently changing the IP addresses of the nodes on a network to neutralize hit list worms; however, this method has many limitations and faces issues including dealing with hosts with static IP addresses or entries in DNS [47]. This is an area for future research.

Only token-based signature detection is able to detect truly polymorphic worms. If the worm fragments itself differently every time it attacks, only systems that are able to handle partial signatures can catch them.

As for metamorphic worms, there are limited solutions available. M. Chouchane and Lakhota proposed the “engine signature” approach to detect metamorphic worms based on a scoring system that would measure the likelihood that the code (worm) might have been generated by a known instruction-substituting metamorphic engine [48].

There has been a lot of work done in worm detection, but there are still more challenges to face in this field.

CONTAINMENT

Detecting worms is important, but it is just as important to stop them from spreading. If a worm can be found ahead of infection, say if the system detects the worm by its signature at the border gateway, the system can try to block the worm and prevent any machine from being infected. But this is not the case most of the time. System administrators and users do not realize there is a worm attack until a victim is having some abnormal behavior and the damage has already been caused. Reacting quickly and minimizing the damage after infection is as important as preventing and detecting worms.

At this point, the worm is found to exist. Those characteristics used for detection no longer matter. We need containment systems to eliminate worms. Many containment methods have been proposed in the past few years. These methods are summarized in Fig. 2 and classified into three categories: slowing down, blocking, and decoying worms.

The first approach is to slow down the spread of worms, and give time for human reaction and intervention. Several methods have been presented, such as using a feedback loop to delay suspicious traffic [49] and using rate limiting techniques at different network levels to slow down infection [50].

These proposals suggest ways to slow down the speed of worm propagation, but worms normally spread at an extremely high speed. Within minutes, worms can spread through whole networks. Human reaction time is a lot slower than a well designed computer system. A good worm containment system should be automated and should not only slow down an infection, but try to stop it.

BLOCKING

Automatically blocking off certain worm-like traffic is another method of containment. When worm-like behavior is discovered, the source has to be isolated from the rest of the network to prevent more machines from getting infected. Blocking is often used together with slowing down. The system can first try to slow down the infection when the first level threshold is met to avoid false positives; if the situation becomes worse and the second threshold is reached, blocking will be utilized.

There are two major approaches to blocking. One is to block off packets with certain content, and the other is to block traffic to and/or from certain addresses. Moore *et al.* [51] simulated containment with both methods, and the result shows that content blocking is more efficient and more effective than address blocking. For either blocking scheme, the challenge is to define when and whom to block to avoid false alarms.

Address Blocking — Address blocking means that when a host is identified as a scanner or victim, any traffic from that host address is dropped. This technique is normally implemented at the border router/gateway, so the containment system is able to perform this task. The system will need to keep a blacklist, which contains addresses to be blocked.

Several algorithms have been proposed for address-block-based containment. One system mentioned before, which uses the success/failure connection ratio for detection, is also designed to block an address when the miss and hit ratio is greater than 10 or another predefined value [30]. The DAW architecture [26] mentioned before also implements address blocking. If certain hosts persistently keep high failure rates, the address is blocked, and the system waits for human intervention for unblocking or further analysis.

Address blocking has to be implemented very carefully to reduce false alarms. In the case of false positives, noninfected hosts might get blocked off, and if the attackers use this loophole, they can trigger a network malfunction, launch DoS attacks, and bring damage to the organization. In addition, Brumley *et al.*'s analysis showed that the effectiveness of address blocking depends on a short reaction time in putting infected hosts on the blacklist. This can especially pose severe challenges in defending against fast propagating worms [52].

Content Blocking — Content blocking is used in most signature-based detection systems. If packet content matches a worm signature, the packet will be dropped automatically. The system can also make decisions based on the type of packet obtained from the header information. Furthermore, certain error messages might not be let through to avoid giving information to a scan source.

In the containment algorithm proposed by Weaver *et al.* [53], after the traffic anomaly reaches a certain threshold, the system will only allow packets from already established connections to go through. Scan-like packets such as TCP SYN to initialize connection and TCP RST not from a pre-established connection will be dropped.

Content blocking allows legitimate traffic to pass while stopping infection-like traffic from going to its destination. It may cause less harm when false alarms arise.

HONEYPOT TO DECOY

HoneyPot was originally designed to lure attackers as a nonexisting host that appears to be valid. Following this spirit, a honeypot can be used as a decoy to lure Internet worms. Worms aggressively find victims to infect. Why not just let them find and infect some fake hosts and leave the real machines alone? Figure 5 shows how a honeypot can be used to contain Internet worms. The top part of the illustration shows that worms infect one machine, then propagate and infect more hosts on the network (shown in solid lines). The bottom part of the figure illustrates the case when worms infect a honeypot machine. The worm stays there without infecting more hosts (shown in dotted lines). This is because the targets the worms have found and try to infect are all simulated by the honeypot, inside which the worms are trapped.

In [39] Provos noted that HoneyD can not only be used in worm detection, but can also be configured to appear as a host having vulnerable applications to decoy and control the spread of worms. HoneyD can emulate a large size of network address space, so it can be used to lure the worm into thinking it is infecting actual hosts while it is actually attacking the honeypot, thus slowing down the infection of real hosts on the network. These honeypots should be installed and activated as soon as a worm is detected. The earlier this activation, the less damage the worm will cause. In the research reported in [39], if the honeypot starts 20 min after the beginning of the worm spread, with a deployment of about 262,000 virtual honeypots, the system is able to stop the worm completely.

FIGHTING WORMS IN DIFFERENT SCOPES

Different detection and containment systems are designed for deployment in different locations of the Internet, which may have different views of network traffic. A system installed at a LAN gateway has the scope of the local network, and can monitor traffic coming in and going out of that network. A system implemented inside an ISP can monitor multiple LANs and has a broader view. A distributed system may gather data across a wider coverage area and monitor all the traffic on the Internet to obtain a global view.

This article discusses worm detection and containment in LAN, ISP, and global scopes, as shown in Fig. 2. These systems can be located inside the network, at a network border, or scattered on the Internet.

LOCATION OF DEFENSE

A common point to place an IDS is at a network boundary such as a gateway or border router. Sitting at the edge of a network, the IDS can inspect all traffic going in and out of the network to discover suspicious packets. This is useful if the system employs signature-based detection algorithms, since the system needs to match the content of each packet with the signature database. It is also good for many anomaly-based

systems because many detection algorithms are based on header information, and several are based on payload information. For example, border routers check the headers of every packet for routing purposes. Some anomaly-based systems can also be located inside the network, depending on the parameters used for detection. For example, honeypots are normally implemented inside a network or at the DMZ and appear to be regular hosts or servers.

A detection algorithm implemented at the border router is good for defending pure random scanning worms and is less likely to catch worms that employ local network preferred scans. Containment is normally better located at the border than inside the network because most of the algorithms try to block or slow down certain traffic, which is easier to execute at the gateway of the network.

In their research on slowing down Internet worms using rate limiting, Wong *et al.* [50] tested rate limiting at individual hosts, edge routers, or backbone routers in simulation. The result shows that rate control at backbone routers is as effective as implementation at all the hosts the backbone routers cover. But it is more complex to install rate limiting filters at every single host than only at the backbone routers. If it is an enterprise environment, they suggest installing rate limiting filters at both the edge router and some portion of the individual hosts to have better protection.

While detection at the host level does not provide many of the benefits of detection at the network border level, such as visibility on all packets going in and out of the network, it provides its own unique benefits and advantages. These advantages include:

- Verification of an attack by checking if the attack or exploit was successful (less prone to false positives when compared to network border level detection); ability to monitor system-specific activities such as adding or removing users, root/admin privileges, and the system logs
- Ability to monitor changes, such as file size or disk space usage, to specific key components of the system such as libraries, DLLs, and executable files
- Ability to monitor and inspect application-specific logs and information not normally available at network level detection systems [54]

However, Brumley *et al.* showed that if local detection and containment is used alone, to be effective there is a need for a very high deployment ratio (to slow down the worm propagation by a factor of two, half of the hosts on the network need to deploy the defense) [52].

Instead of having the detection or containment system at one single location, distributed sensors and containment systems also have their advantage. A distributed system can cover a greater portion of the network and have a broader view, and be able to stop the worms faster and more effectively. Malan and Smith [55] proposed a collaborative detection system to reduce the false positive in host-based anomaly detection systems by defining a host as behaving anomalously if its behavior correlates too well with other networked but independent hosts. Using this system, they were able to distinguish non-worm processes on a system from worm processes 99 percent of the time [55].

Stolfo [56] proposed the Worminator project that detects, reports, and defends against early attack events. This ongoing project is a collaboration of several academic institutions including Columbia University, GIT, FIT, MIT, and Syracuse University. They use Antura sensors and the Columbia PAYL sensors [57] to detect stealthy scans and probes from outside of the firewalls. Analyzing the feedback of these sensors provides a greater picture of worm behavior.

In [51] Moore *et al.* simulated Internet worm containment in different percentages of customer autonomous systems (ASs) and different coverage of top ISP ASs. The result shows that implementing in the ISP ASs is far more effective than in the customers' ASs. In order for the containment to be useful, the paths covered by the top 100 largest ASs have to be included. This means almost all Internet paths have to employ the containment system and requires wide cooperation among ISPs, which is very difficult to achieve.


SCOPE OF DEFENSE

Different locations of implementation give different scopes of worm detection and containment. Figure 6 illustrates the correlation of these different network levels. Most IDSs are designed for local area or enterprise network detection and containment. The local area or enterprise network is a clearly defined entity and is normally controlled by one single organization, which has central management when it comes to deciding the type of IDS to implement.

A broader extent that is well defined, is the scope of an ISP or AS, which has multiple customer networks connected to it. After detecting worms from a certain customer network, the ISP can slow down or block off partial traffic from that network to prevent worms from spreading to other customer networks. The detection might be more complex for signature-based algorithms because of dealing with large amounts of traffic, so anomaly-based algorithms may be more feasible. Wagner and Plattner proposed an entropy²-based anomaly detection system to detect worms in fast IP networks (networks with large amounts of traffic) such as the Internet backbones [58]. Essentially, the larger the coverage, the more accurate the normal model definition. The DAW architecture [26] is another example of a system for this scope, which is implemented with anomaly-based detection and containment inside an ISP network.

Different parameters are used for different scopes of detection. As the scope grows bigger, detection may be rougher, but containment is more effective. Worms spread at very fast speeds. The damage of a worm outbreak is normally very broad, often across countries. In previous sections we have seen that many of the worm detection algorithms are implemented based on monitoring larger size networks (220 nodes or more). Moore *et al.* [51] also showed that worm containment is only practical if a large fraction of the Internet unites and works together. This leads us to conclude that global scope is necessary in defending against worms. Zhang *et al.* [11] proposed a system combining both control plane data (routing data) as well as data plane data (packet headers and payloads) to detect and contain Internet worms more effectively. In this system, anomalies detected on the data plane are used to identify ASs that are associated with the attacks and apply control plane filters to contain them. Furthermore, anomalies detected on the control plane (e.g., IP hijacking) can be used to deploy strict data plane controls on a particular AS [11].

The idea of setting up a Center for Disease Control (CDC) for global scope detection was brought up by Staniford, Pax-



■ Figure 6. Different scopes of detection and containment.

son, and Weaver [5]. They believe that the CDC should have the following roles: identifying outbreaks, rapidly analyzing pathogens, fighting infections, anticipating new vectors, proactively devising detectors for new vectors, and resisting future threats. This center should be deployed across the globe. There are real benefits in this approach when “one’s allies are awake and working while one sleeps.”

Qin *et al.* [32] suggested that a CDC is not very practical if used by itself. One reason is privacy: not all organizations are willing to share their data with others. Another reason stated in the article is that the architecture of a CDC requires a victim: a participant will not hear about the worm outbreak until there is a victim, and this victim could be any participant.

The detection system utilizing ICMP error messages [25] discussed earlier is another system that tries to obtain global scope. This global-scale worm detection and analysis system is based on ICMP destination unreachable error messages forwarded by routers. This method is not possible until there are enough participating routers. To deal with such large amounts of data, this central point will need a considerable amount of resources for processing. The global scope is essential, but very hard to achieve, mainly because it requires wide cooperation among ISPs, organizations, and countries, and there are also privacy issues. In addition, these entities might have conflicts of interest that can make cooperation very difficult.

Detection and containment have to be implemented with a hierarchical approach. Most enterprises already have their own security systems to protect their networks. Other LANs should do the same to defend attacks at the lowest level, even going as far as having detection mechanisms on some individual hosts as discussed in the last section. Detections at the LAN level are more detailed than at the ISP level, and if these local networks can flag the ISPs when worms are found, ISPs can then confirm this situation and start containment procedures. ISPs have a great position in worm containment since they are the junction of data exchange. At the same time, an ISP can alert other ISPs of its findings about a worm and take necessary precautions. Since a worm outbreak can be a worldwide disaster, everybody should work together. Briesemeister and Porras [59] presented an approach to evaluate collaborative worm defense mechanisms against future, unseen, and possibly defense-aware worms. In their model Briesemeister and Porras do not assume any specific worm propagation strategy and consider all possible infection sequences, and propose that studying these propagation

² Entropy is a measure on how random the traffic is and comes into perspective as worm traffic is more structured in some respects and more random in others when compared to normal network traffic [58].

sequences will result in understanding how current worm defense algorithms can be improved to prevent worms with similar patterns from succeeding in the future [59].

CONCLUSION

We have identified the characteristics of existing and hypothetical worms during the target finding and propagation phases of a worm's life cycle. They are classified based on target finding, propagation, transmission scheme, and payload format. Current detection algorithms are organized based on the categories of signature-based, anomaly-based, or hybrid. We have evaluated these categories against worm characteristics. We have classified current containment schemes based on the methods they use to control the spread of worms. We have also explored the implementations of detection and containment at different network locations and system scopes.

An ideal system should use a combination of schemes to have more comprehensive coverage. Different detection schemes are useful at different levels of implementation. So far, there is no ultimate solution to deal with all existing and hypothetical worms. New attack technologies are being developed every day, and the threat constantly exists. We have pointed out the remaining challenges and future work to be done based on the analysis of current algorithms. So far, there are limited solutions for detecting passive and topological scanning worms, flash worms, and metamorphic worms; nevertheless, as pointed out in research by Kienzle and Elder [60], the majority of new worms coming out every day are not novel and are derivative in nature. As a result, by defending against yesterday's worms, we can effectively protect ourselves against most new worms; at the same time, we also need to prepare for the threats of new novel worms that can hit us in the future [60].

REFERENCES

- [1] E. Spafford, "The Internet Worm Program: An Analysis," *Comp. Commun. Rev.*, 1989.
- [2] D. T. C. Zou, W. Gong, and S. Cai, "Routing Worm: A Fast, Selective Attack Worm Based on IP Address Information," *Proc. 19th ACM/IEEE/SCS Wksp. Principles of Advanced and Distrib. Simulation*, 2005.
- [3] J. Z. Chen, "A Self-Learning Worm Using Importance Scanning," *Proc. ACM CCS Wksp. Rapid Malcode (WORM '05)*, 2005.
- [4] F. M. Moheeb Abu Rajab and Andreas Terzis, "On the Impact of Dynamic Addressing on Malware Propagation," *Proc. ACM WORM '06*, 2006.
- [5] V. P. S. Staniford and N. Weaver, "How to Own the Internet in Your Spare Time," *Proc. 11th USENIX Sec. Symp.*, 2002.
- [6] D. M. S. Staniford, V. Paxson, and N. Weaver, "The Top Speed of Flash Worms," *Proc. ACM WORM '04*, 2004.
- [7] V. P. N. Weaver, S. Staniford, and R. Cunningham, "A Taxonomy of Computer Worms," *Proc. ACM WORM '03*, 2003.
- [8] J. M. Niels Provos and Ke Wang, "Search Worms," *Proc. ACM WORM '06*, 2006.
- [9] G. P. Schaffer, "Worms and Viruses and Botnets, Oh My! Rational Responses to Emerging Internet Threats," *IEEE Sec. & Privacy*, vol. 4, 2006, pp. 52–58.
- [10] J. Z. Moheeb Abu Rajab, F. Monrose, and A. Terzis, "A Multifaceted Approach to Understanding the Botnet Phenomenon," *Proc. 6th ACM SIGCOMM on Internet Measurement*, 2006.
- [11] E. C. Ying Zhang and Z. Morley Mao, "Internet-Scale Malware Mitigation: Combining Intelligence of the Control and Data Plane," *Proc. ACM WORM '06*, 2006.
- [12] V. P. D. Moore et al., "Inside the Slammer Worm," *IEEE Sec. & Privacy*, vol. 1, 2003, pp. 33–39.
- [13] D. G. Glazer, "Computer Worms," May 2005, <http://www.research.umbc.edu/~dgorin1/is432/worms.htm>
- [14] "Morris (Computer Worm)," retrieved July 2007, http://en.wikipedia.org/wiki/Morris_worm
- [15] "F-Secure Virus Descriptions: Nimda," retrieved July 2007, <http://www.f-secure.com/v-descs/nimda.shtml>, 2001
- [16] "CERT" Advisory CA-2001-26 Nimda Worm," retrieved July 2007, <http://www.cert.org/advisories/CA-2001-26.html>, 2001.
- [17] "F-Secure Computer Virus Information Pages: Slammer," vol. May, 2005, <http://www.f-secure.com/v-descs/mssqlm.shtml>
- [18] "Sasser Worm Analysis — LURHQ," May 2005, <http://www.lurhq.com/sasser.html>
- [19] "Secunia — Virus Information — Sasser.G," May 2005, http://secunia.com/virus_information/11515/sasser.g
- [20] "F-Secure Computer Virus Information Pages: Witty," May 2005, <http://www.f-secure.com/v-descs/witty.shtml>
- [21] C. S. D. Moore, "The Spread of the Witty Worm," *IEEE Sec. & Privacy*, vol. 2, 2004, pp. 46–50.
- [22] R. A. et al., *Snort 2.1 Intrusion Detection*, 2nd ed., Syngress, O'Reilly, 2004, pp. 490–91.
- [23] S. E. D. Bolzoni and P. Hartel, "POSEIDON: A 2-Tier Anomaly-Based Network Intrusion Detection System," *Proc. 4th IEEE Int'l Wksp. Info. Assurance*, 2006.
- [24] "Snort," May 2005, <http://www.snort.org>
- [25] G. B. V. Berk and R. Morris, "Designing a Framework for Active Worm Detection on Global Networks," *Proc. 1st IEEE Int'l Wksp. Info. Assurance*, 2003.
- [26] S. C. Y. Tang, "Slowing Down Internet Worms," *Proc. 24th IEEE Int'l Conf. Distrib. Comp. Sys.*, 2004.
- [27] S. C. S. Ranka, "An Internet-Worm Early Warning System," *Proc. IEEE GLOBECOM*, 2004.
- [28] S. S. J. Jung, A. Berger, "Fast Detection of Scanning Worm Infections," *Proc. 7th Int'l Symp. Recent Advances in Intrusion Detection*, 2004.
- [29] V. P. J. Jung et al., "Fast Portscan Detection Using Sequential Hypothesis Testing," *Proc. IEEE Symp. Sec. and Privacy*, 2004.
- [30] D. E. N. Weaver, S. Staniford, and V. Paxson, "Worms vs. Perimeters — The Case for Hard-LANs," *Proc. 12th Annual IEEE Symp. High Perf. Interconnects*, 2004.
- [31] M. S. G. Gu et al., "Worm Detection, Early Warning and Response Based on Local Victim Information," *Proc. 20th Annual Comp. Sec. Apps. Conf.*, 2004.
- [32] D. D. X. Qin et al., "Worm Detection Using Local Networks," tech. rep., College of Computing, Georgia Tech, 2004.
- [33] C. E. S. Singh, G. Varghese, and S. Savage, "The EarlyBird System for Real-Time Detection of Unknown Worms," UCSD tech rep. CS2003-0761, 2003.
- [34] L. G. Z. Chen and K. Kwiat, "Modeling the Spread of Active Worms," *Proc. IEEE Comp. and Commun. Societies Annual Joint Conf.*, 2003.
- [35] S. V. J. Wu, L. Gao, and K. Kwiat, "An Efficient Architecture and Algorithm for Detecting Worms with Various Scan Techniques," *Proc. Network and Distrib. Sys. Sec. Symp.*, 2004.
- [36] L. Spitzner, *Honeypot: Tracking Hackers*, Addison-Wesley, 2002.
- [37] X. Q. D. Dagon et al., "Honeystat: Local Worm Detection Using Honeypots," *Proc. 7th Symp. Recent Advances in Intrusion Detection*, 2004.
- [38] "Honeyd Virtual Honeypot," retrieved Aug. 2005, <http://honeyd.org>
- [39] N. Provos, "A Virtual Honeypot Framework," *Proc. 13th USENIX Sec. Symp.*, 2004.
- [40] P. K. C. Matthew V. Mahoney, "PHAD: Packet Header Anomaly Detection for Identifying Hostile Network Traffic," FL Inst. Tech. tech. rep. CS-2001-04, 2001.
- [41] S. J. S. Ke Wang, "Anomalous Payload-Based Network Intrusion Detection," *Proc. Recent Advances in Intrusion Detection*, 2004.
- [42] G. C. Ke Wang and Salvatore J. Stolfo, "Anomalous Payload-based Worm Detection and Signature Generation," *Proc. 8th Int'l Symp. Recent Advances in Intrusion Detection*, 2005.
- [43] B. M. J. Lockwood, "Design of a System for Real-Time Worm Detection," *Proc. 12th IEEE Annual Symp. High Perf. Interconnects*, 2004.
- [44] J. C. C. Kreibich, "Honeycomb — Creating Intrusion Detection Signatures Using Honeypots," *Proc. 2nd Wksp. Hot Topics in Networks*, 2003.
- [45] B. K. H. Kim, "Autograph: Toward Automated, Distributed Worm Signature Detection," *Proc. 13th USENIX Sec. Symp.*, 2004.
- [46] B. K. J. Newsome and D. Song, "Polygraph: Automatically

- Generating Signatures for Polymorphic Worms," *Proc. IEEE Symp. Sec. and Privacy*, 2005.
- [47] P. A. S. Antonatos, E. P. Markatos, K. G. Anagnostakis, "Defending Against Hitlist Worms Using Network Address Space Randomization," *Proc. ACM WORM '05*, 2005.
- [48] A. L. Mohamed R. Chouchane, "Using Engine Signature to Detect Metamorphic Malware," *Proc. ACM WORM '06*, 2006.
- [49] J. C. R. Dantu and A. Yelimeli, "Dynamic Control of Worm Propagation," *Proc. Int'l Conf. Info. Technology: Coding and Comp.*, 2004.
- [50] C. W. C. Wong et al., "Dynamic Quarantine of Internet Worms," *Proc. Int'l Conf. Dependable Sys. and Networks*, 2004.
- [51] C. S. D. Moore, G. M. Voelker, and S. Savage, "Internet Quarantine: Requirements for Containing Self-Propagating Code," *Proc. IEEE INFOCOM*, 2003.
- [52] L.-H. L. David Brumley, Pongsin Poosankam, and Dawn Song, "Design Space and Analysis of Worm Defense Strategies," *Proc. ACM Symp. Info., Comp. and Commun. Sec.*, 2006.
- [53] S. S. N. Weaver and V. Paxson, "Very Fast Containment of Scanning Worms," *Proc. 13th USENIX Sec. Symp.*, 2004.
- [54] B. Laing, "How To Guide-Implementing a Network Based Intrusion Detection System," retrieved July 2007, <http://www.snort.org/docs/#deploy>
- [55] D. J. M. a. M. D. Smith, "Exploiting Temporal Consistency to Reduce False Positives in Host-Based, Collaborative Detection of Worms," *Proc. ACM WORM '06*, 2006.
- [56] S. J. Stolfo, "Worm and Attack Early Warning," *IEEE Sec. & Privacy*, vol. 2, 2004, pp. 73–75.
- [57] "Worminator Projector," retrieved Apr., 2005, <http://worminator.cs.columbia.edu/public/index.jsp>
- [58] B. P. A. Wagner, "Entropy Based Worm and Anomaly Detection in Fast IP Networks," *Proc. 14th IEEE WET ICE/STCA Sec. Wksp.*, 2005.
- [59] P. A. P. L. Briesemeister, "Automatically deducing Propagation Sequences that Circumvent a Collaborative Worm Defense," *Proc. 25th IEEE Conf. Perf., Comp., and Commun.*, 2006.
- [60] M. E. D. Kienzle, "Recent Worms: A Survey and Trends," *Proc. ACM WORM '03*, 2003.

BIOGRAPHIES

PELE LI was born in Taipei, Taiwan. She received her M.S. degree in computer engineering from San Jose State University, California, where she gained in-depth knowledge of intrusion detection systems and flash worms from her research and study. She has been involved in network security since 1999 and currently consults on IT management projects.

MEHDI SALOUR received his M.S. degree in software engineering from San Jose State University with research concentrations in networking software and security. He is currently serving as vice president of service delivery and support at 8x8, Inc. where he oversees the entire customer experience from production network operations to customer support for the company's voice and video over IP services.

XIAO SU (xsu@email.sjsu.edu) received her Ph.D. in computer science from the University of Illinois at Urbana-Champaign. She is currently an associate professor in the Computer Engineering Department, San Jose State University. Her research interests include network security, multimedia communications, media coding, and mobile computing. She is a recipient of the National Science Foundation CAREER award.